

Notes on Jumping PRNGs Ahead

This version of the document is dated 2023-02-10.

Peter Occil

Some pseudorandom number generators (PRNGs) have an efficient way to advance their state as though a huge number of PRNG outputs were discarded. Notes on how they work are described in the following sections.

0.1 F_2 -linear PRNGs

For some PRNGs, each bit of the PRNG's state can be described as a linear recurrence on its entire state. These PRNGs are called *F_2 -linear PRNGs*, and they include the following:

- Linear congruential generators (LCGs) with a power-of-two modulus.
- Xorshift PRNGs.
- PRNGs in the xoroshiro and xoshiro families.
- Linear or generalized feedback shift register generators, including Mersenne Twister.

For an F_2 -linear PRNG, there is an efficient way to discard a given (and arbitrary) number of its outputs (to "jump the PRNG ahead"). This jump-ahead strategy is further described in (Haramoto et al., 2008)¹. See also (Vigna 2017)². To calculate the jump-ahead parameters needed to advance the PRNG N steps:

1. Build M , an $S \times S$ matrix of zeros and ones that describes the linear transformation of the PRNG's state, where S is the size of that state in bits. For an example, see sections 3.1 and 3.2 of (Blackman and Vigna 2019)³, where it should be noted that the additions inside the matrix are actually XORs.
2. Find the *characteristic polynomial* of M . This has to be done in the two-element field F_2 , so that each coefficient of the polynomial is either 0 or 1.

For example, SymPy's `charpoly()` method alone is inadequate for this purpose, since it doesn't operate on the correct field. However, it's easy to adapt that method's output for the field F_2 : even coefficients become zeros and odd coefficients become ones.

Note that for a linear feedback shift register (LFSR) generator, the characteristic polynomial's coefficients are 1 for each of its "taps" (and "tap" 0), and 0 elsewhere. For example, an LFSR generator with taps 6 and 8 has the characteristic polynomial $x^8 + x^6 + 1$.

The section "Jump Parameters for Some PRNGs" shows characteristic polynomials for some PRNGs and one way their coefficients can be represented.

3. Calculate `powmodf2(2, N, CP)`, where `powmodf2` is a modular power function that calculates $2^N \bmod CP$ in the field F_2 , and `CP` is the characteristic polynomial. (N is the number of PRNG outputs to discard.) Regular modular power functions, such as `BigInteger's modPow` method, won't work here, even if the polynomial is represented in the manner described in "Jump Parameters for Some PRNGs".
4. The result is a *jump polynomial* for jumping the PRNG ahead N steps, that is, for discarding N outputs of the PRNG.

An example of its use is found in the `jump` and `long_jump` functions in the [xoroshiro128plus source code](#), which are identical except for the jump polynomial. In both functions, the jump polynomial's coefficients are packed into a 128-bit integer (as described in "Jump Parameters for Some PRNGs"), which is then split into the lower 64 bits and the upper 64 bits, in that order.

0.2 Counter-Based PRNGs

Counter-based PRNGs, in which their state is updated simply by incrementing a counter, can be trivially jumped ahead just by changing the seed, the counter, or both (Salmon et al. 2011)⁴.

0.3 Multiple Recursive Generators

A *multiple recursive generator* (MRG) generates numbers by transforming its state using the following formula: $x(k) = (x(k-1)*A(1)$

$+ x(k-2)*A(2) + \dots + x(k-n)*A(n)) \bmod \text{modulus}$, where $A(i)$ are the *multipliers* and modulus is the *modulus*.

For an MRG, the following matrix (M) describes the state transition $[x(k-n), \dots, x(k-1)]$ to $[x(k-n+1), \dots, x(k)] \pmod{\text{modulus}}$:

$$\begin{array}{c|ccccc|} 0 & 1 & 0 & \dots & 0 & \\ 0 & 0 & 1 & \dots & 0 & \\ \cdot & \cdot & \cdot & \dots & \cdot & \\ 0 & 0 & 0 & \dots & 1 & \\ A(n)A(n-1) & A(n-1) & A(n-2) & \dots & A(1) & \\ \hline & -1 & -2 & & & \end{array}$$

To calculate the parameter needed to jump the MRG ahead N steps, calculate $M^N \bmod \text{modulus}$; the result is a *jump matrix* J .

Then, to jump the MRG ahead N steps, calculate $J * S \bmod \text{modulus}$, where J is the jump matrix and S is the state in the form of a column vector; the result is a new state for the MRG.

This technique was mentioned (but for binary matrices) in Haramoto, in sections 1 and 3.1. They point out, though, that it isn't efficient if the transition matrix is large. See also (L'Ecuyer et al., 2002)⁵.

0.3.1 Example

A multiple recursive generator with a modulus of 1449 has the following transition matrix:

$$\begin{array}{c|ccc|} 0 & 1 & 0 & & \\ 0 & 0 & 1 & & \\ 444 & 342 & 499 & & \end{array}$$

To calculate the 3×3 jump matrix to jump 100 steps from this MRG, raise this matrix to the power of 100 then take the result's elements $\bmod 1449$. One way to do this is the "square-and-multiply" method, described by D. Knuth in *The Art of Computer Programming*: Set J to the identity matrix, N to 100, and M to a copy of the transition matrix, then while N is greater than 0:

1. If N is odd⁶, multiply J by M then take J 's elements $\bmod 1449$.
2. Divide N by 2 and round down, then multiply M by M then take M 's elements $\bmod 1449$.

The resulting J is a *jump matrix* as follows:

	156	93	1240	
	1389	1128	130	
	1209	930	793	

Transforming the MRG's state with J (and taking its elements mod 1449) will transform the state as though 100 outputs were discarded from the MRG.

0.4 Linear Congruential Generators

A *linear congruential generator* (LCG) generates numbers by transforming its state using the following formula: $x(k) = (x(k-1)*a + c) \bmod \text{modulus}$, where a is the *multiplier*, c is the additive constant, and modulus is the *modulus*.

An efficient way to jump an LCG ahead is described in (Brown 1994)⁷. This also applies to LCGs that transform each $x(k)$ before outputting it, such as M.O'Neill's PCG32 and PCG64.

An MRG with only one multiplier expresses the special case of an LCG with $c = 0$ (also known as a *multiplicative* LCG). For c other than 0, the following matrix describes the state transition $[x(k-1), 1]$ to $[x(k), 1] \pmod{\text{modulus}}$:

	a	c	
	0	1	

Jumping the LCG ahead can then be done using this matrix as described in the previous section.

0.5 Multiply-with-Carry, Add-with-Carry, Subtract-with-Borrow

There are implementations for jumping a multiply-with-carry (MWC) PRNG ahead, but **only in source-code form**. I am not aware of an article or paper that describes how jumping an MWC PRNG ahead works.

I am not aware of any efficient ways to jump an add-with-carry or subtract-with-borrow PRNG ahead an arbitrary number of steps.

0.6 Combined PRNGs

A combined PRNG can be jumped ahead N steps by jumping each of its components ahead N steps.

0.7 Jump Parameters for Some PRNGs

The following table shows the characteristic polynomial and jump polynomials for some PRNG families. In the table:

- Each number before the colon in the jump polynomial column is the number of PRNG outputs discarded when the corresponding jump polynomial is used.
- Each polynomial's coefficients are zeros and ones, so the table shows them as a base-16 integer that stores the coefficients as individual bits: the least significant bit is the degree-0 coefficient, the next bit is the degree-1 coefficient, and so on. For example, the integer 0x23 stores the coefficients of the polynomial $x^5 + x + 1$.
- Each characteristic polynomial's highest coefficient is x^n , where n is the PRNG's state size. Thus, the table shows it as a base-16 integer with n plus one bits.
- "'Period'/φ" means the PRNG's maximum cycle length divided by the golden ratio, and rounded to the closest odd integer; this jump parameter is chosen to avoid overlapping number sequences as much as possible (see also [NumPy documentation](#)).

PRNG	Characteristic Polynomial	
xoroshiro64	0x1053be9da6e2286c1	2 ³² : 0x
		2 ⁴⁸ : 0x
		"Period"
xoshiro128	0x100fc65a2006254b11b489db6de18fc01	2 ³² : 0x
		2 ⁴⁸ : 0x
		2 ⁶⁴ : 0x
		2 ⁹⁶ : 0x
		"Period"
		2 ³² : 0x
		2 ⁴⁸ : 0x

xoroshiro128 (except ++)	0x10008828e513b43d5095b8f76579aa001	2 ⁶⁴ : 0x
		2 ⁹⁶ : 0x
xoroshiro128++	0x10031bcf2f855d6e58dae70779760b081	"Period
		2 ³² : 0x
		2 ⁴⁸ : 0x
		2 ⁶⁴ : 0x
		2 ⁹⁶ : 0x
		"Period
		2 ³² :
		0xe05f
		2 ⁴⁸ :
		0x5f72
		2 ⁶⁴ :
		0x12e4
		2 ⁹⁶ :
		0x31e4
xoshiro256	0x10003c03c3f3ecb1904b4edcf26259f85- 0280002bcefd1a5e9d116f2bb0f0f001	2 ¹²⁸ :
		0x39a1
		2 ¹⁶⁰ :
		0xf567
		2 ¹⁹² :
		0x3910
		2 ²²⁴ :
		0xa2b1
		"Period
		0x2944

0.8 Acknowledgments

Sebastiano Vigna reviewed this page and gave comments.

1 Notes

1. Haramoto, Matsumoto, Nishimura, Panneton, L'Ecuyer, "Efficient Jump Ahead for F_2 -Linear Random Number Generators", *INFORMS Journal on Computing* 20(3), Summer 2008.↵
2. Vigna, S., "Further scramblings of Marsaglia's xorshift generators", *Journal of Computational and Applied Mathematics* 315 (2017).↵
3. Blackman, Vigna, "Scrambled Linear Pseudorandom Number Generators", 2019.↵
4. Salmon, John K., Mark A. Moraes, Ron O. Dror, and David E. Shaw. "Parallel random numbers: as easy as 1, 2, 3." In *Proceedings of 2011 International Conference for High Performance Computing, Networking, Storage and Analysis*, pp. 1-12. 2011.↵
5. L'Ecuyer, Simard, Chen, Kelton, "An Object-Oriented Random-Number Package with Many Long Streams and Substreams", *Operations Research* 50(6), 2002.↵
6. "x is odd" means that x is an integer and not divisible by 2. This is true if $x - 2*\text{floor}(x/2)$ equals 1, or if x is an integer and the least significant bit of abs(x) is 1.↵
7. Brown, F., "Random Number Generation with Arbitrary Strides", *Transactions of the American Nuclear Society* Nov. 1994.↵